

PACKAGE CRS FAQ

JEFFREY S. RACINE

CONTENTS

1. Overview and Current Version	2
2. Frequently Asked Questions	2
2.1. How do I cite the <code>crs</code> package?	2
2.2. I have never used R before. Can you direct me to some introductory material that will guide me through the basics?	3
2.3. How do I keep all R packages on my system current?	3
2.4. It seems that there are a lot of packages that must be installed in order to conduct econometric analysis (<code>tseries</code> , <code>lmtest</code> , <code>np</code> , etc.). Is there a way to avoid having to individually install each package individually?	4
2.5. Is there a ‘gentle guide’ to the <code>crs</code> package that contains some easy to follow examples?	4
2.6. I noticed you have placed a new version of the <code>crs</code> package on CRAN. How can I determine what has been changed, modified, fixed etc?	4
2.7. How can I read data stored in various formats such as Stata, SAS, Minitab, SPSS etc. into the R program?	4
2.8. Where can I get some examples of R code for the <code>crs</code> package in addition to the examples in the help files?	5
2.9. Can I use <code>crs</code> to perform linear regression?	5
2.10. I would like more/less information displayed when conducting search using the NOMAD routines...	6
2.11. When I have a large number of regressors/data the function <code>crs</code> just ‘sits there’ when conducting cross-validation via NOMAD...	6
2.12. My estimated model is not ‘smooth’ (e.g. cross-validation chooses, say, the spline degree=1 and number of segments=3). How can I modify this?	7
2.13. I estimated a parametric model using the <code>lm()</code> function. How can I compare the cross-validation score from the <code>crs()</code> approach with that for the parametric model?	7
2.14. Why do some runs result in a function value of 1.340781e+154 when conducting multistarting?	8
2.15. <code>snomadr</code> appears to be crashing	8
2.16. How can I save a PDF of a plot created with the option <code>persp.rgl=TRUE</code> ?	8

Date: December 3, 2011.

2.17. I have noticed that as more categorical predictors are added or as the number of outcomes for each categorical predictor increases, computation time increases when <code>kernel=TRUE</code> . Why is this so and can anything be done?	9
References	10
Changes from Version 0.15-9 to 0.15-10 [3-Dec-2011]	11
Changes from Version 0.15-8 to 0.15-9 [25-Nov-2011]	11
Changes from Version 0.15-7 to 0.15-8 [16-Nov-2011]	11
Changes from Version 0.15-6 to 0.15-7 [24-Oct-2011]	12
Changes from Version 0.15-5 to 0.15-6 [17-Oct-2011]	12
Changes from Version 0.15-4 to 0.15-5 [16-Oct-2011]	12
Changes from Version 0.15-3 to 0.15-4 [15-Oct-2011]	12
Changes from Version 0.15-2 to 0.15-3 [05-Sept-2011]	12
Changes from Version 0.15-1 to 0.15-2 [30-July-2011]	13
Changes from Version 0.15-0 to 0.15-1 [29-Jul-2011]	13
Changes from Version 0.14-9 to 0.15-0 [23-Jun-2011]	14
Changes from Version 0.14-8 to 0.14-9 [20-Jun-2011]	14
Changes from Version 0.14-7 to 0.14-8 [10-Jun-2011]	14
Version 0.14-7 [09-Jun-2011]	14

1. OVERVIEW AND CURRENT VERSION

This set of frequently asked questions is intended to help users who are encountering unexpected or undesired behavior when trying to use the `crs` package.

Kindly report any issues you encounter to me, and please include your code, data, version of the package and version of R used so that I can help track down any such issues (racinej@mcmaster.ca). And, of course, if you encounter an issue that you think might be of interest to others, kindly email me the relevant information and I will incorporate it into this FAQ.

This FAQ refers to the most recent version, which as of this writing is 0.15-10. Kindly update your version should you not be using the most current (from within R, `update.packages()` ought to do it, though also see 2.3 below.). See the appendix in this file for cumulative changes between this and previous versions of the `crs` package.

2. FREQUENTLY ASKED QUESTIONS

2.1. How do I cite the `crs` package? Once you have installed the `crs` package (`install.packages("crs")`), if you load the `crs` package (`library("crs")`) and type `citation("crs")` you will be presented with the following information.

```
> citation("crs")
```

To cite package `crs` in publications use:

Jeffrey S. Racine <racinej@mcmaster.ca> and Zhenghua Nie
<niez@mcmaster.ca> (2011). `crs`: Categorical Regression Splines. R
package version 0.15-10.

A BibTeX entry for LaTeX users is

```
@Manual{,
  title = {crs: Categorical Regression Splines},
  author = {Jeffrey S. Racine and Zhenghua Nie},
  year = {2011},
  note = {R package version 0.15-10},
}
```

2.2. I have never used R before. Can you direct me to some introductory material that will guide me through the basics? There are many excellent introductions to the R environment with more on the way. First, I would recommend going directly to the R website (<http://www.r-project.org>) and looking under Documentation/Manuals (<http://cran.r-project.org/manuals.html>) where you will discover a wealth of documentation for R users of all levels. See also the R task views summary page (<http://cran.nedmirror.nl/web/views/index.html>) for information grouped under field of interest. A few documents that I mention to my students which are tailored to econometricians include <http://cran.r-project.org/doc/contrib/Verzani-SimpleR.pdf>, Cribari-Neto & Zarkos (1999) [1], Racine & Hyndman (2002) [5] and Farnsworth (2006) [3], to name but a few.

Those looking for exemplar data sets outside of those contained in the `crs` package are directed to the `Ecdat` [2] and `AER` [4] packages.

Often the best resource is right down the hall. Ask a colleague whether they use or know anyone who uses R, then offer to buy that person a coffee and along the way drop something like “I keep hearing about the R project... I feel like such a Luddite...”

2.3. How do I keep all R packages on my system current? Run the command `update.packages(checkBuilt=TRUE,ask=FALSE)`, which will not only update all packages that are no longer current, but will also update all packages built under outdated installed versions of R, if appropriate.

2.4. It seems that there are a lot of packages that must be installed in order to conduct econometric analysis (`tseries`, `lmtest`, `np`, etc.). Is there a way to avoid

having to individually install each package individually? Certainly. The Comprehensive R Archive Network (CRAN) is a network of ftp and web servers around the world that store identical, up-to-date, versions of code and documentation for R. The CRAN task view for computational econometrics might be of particular interest to econometricians. The econometric task view provides an excellent summary of both parametric and nonparametric econometric packages that exist for the R environment and provides one-stop installation for these packages.

See cran.r-project.org/web/views/Econometrics.html for further information.

To automatically install a task view, the `ctv` package first needs to be installed and loaded, i.e.,

```
install.packages("ctv")
library("ctv")
```

The econometric task view can then be installed via `install.views()` and updated via `update.views()` (which first assesses which of the packages are already installed and up-to-date), i.e.,

```
install.views("Econometrics")
```

or

```
update.views("Econometrics")
```

2.5. Is there a ‘gentle guide’ to the `crs` package that contains some easy to follow examples? Perhaps the most gentle introduction is contained in the `crs` package itself in the form of a ‘vignette’. To view the vignette run R, install the `crs` package (`install.packages("crs")`), then type `vignette("crs",package="crs")` to view or print the vignette.

See also `vignette("spline_primer",package="crs")` for a vignette that presents a ‘gentle’ introduction to regression splines.

For a listing of all routines in the `crs` package type: `library(help="crs")`.

2.6. I noticed you have placed a new version of the `crs` package on CRAN. How can I determine what has been changed, modified, fixed etc? See the CHANGELOG on the CRAN site (<http://cran.r-project.org/web/packages/crs/ChangeLog>), or go to the end of this document where the CHANGELOG is provided for your convenience.

2.7. How can I read data stored in various formats such as Stata, SAS, Minitab, SPSS etc. into the R program? Install the foreign library via `install.packages("foreign")` then do something like

```
mydat <- read.dta("datafile.dta"),
```

where `datafile.dta` is the name of your Stata data file. Note that, as of version 0.8-34, the foreign package function `read.dta` supports reading files directly over the Internet making for more portable code. For instance, one could do something like

```
mydat <- read.dta(file="http://www.principlesofeconometrics.com/stata/mroz.dta")
```

as one could always do with, say, `read.table()`.

2.8. Where can I get some examples of R code for the crs package in addition to the examples in the help files? Start R then type `demo(package="crs")` and you will be presented with a list of demos for constrained estimation, inference, and so forth. To run one of these demos type, for example, `demo(radial_rgl)` (note that you must first install the `rgl` package to run this particular demo).

To find the location of a demo type `system.file("demo", "radial_rgl.R", package="crs")` for example, then you can take the source code for this demo and modify it for your particular application.

2.9. Can I use crs to perform linear regression? Certainly! Simply use the options `cv="none"`, `degree=rep(1,q)` and `segments=rep(1,q)` where `q` is the number of continuous predictors and `kernel=FALSE` as per the following example:

```
n <- 1000
x1 <- runif(n)
x2 <- runif(n)
z1 <- rbinom(n,1,.1)
z2 <- rbinom(n,1,.1)
y <- x1+x2^2+z1+z2+rnorm(n)
z1 <- factor(z1)
z2 <- factor(z2)
model.lm <- lm(y~x1+x2+z1+z2)
summary(model.lm)
model.crs <- crs(y~x1+x2+z1+z2,cv="none",kernel=FALSE,degree=rep(1,2),segments=rep(1,2))
summary(model.crs)
```

You will note that the summary statistics are identical for each model. Also, when conducting search the initial values are set so that the first model estimated is linear (albeit with `kernel=TRUE` by default) so if the linear model is optimal it will be the one chosen by cross-validation in probability.

2.10. I would like more/less information displayed when conducting search using the NOMAD routines... This is accomplished by feeding the argument `opts=list("DISPLAY_DEGREE"=x)` to `crs` where x is a non-negative integer. Setting $x=0$ produces no information whatsoever while integers $x \geq 1$ provide successively more information.

2.11. When I have a large number of regressors/data the function `crs` just ‘sits there’ when conducting cross-validation via NOMAD... First, if you are concerned that the code is indeed just ‘sitting there’, you can verify that search is progressing by changing the "DISPLAY_DEGREE" setting in the `opts` list along the lines of the following:

```
opts <- list("MAX_BB_EVAL"=10000,
            "EPSILON"=.Machine$double.eps,
            "INITIAL_MESH_SIZE"="r1.0e-01",
            "MIN_MESH_SIZE"=paste("r",sqrt(.Machine$double.eps),sep=""),
            "MIN_POLL_SIZE"=paste("r",sqrt(.Machine$double.eps),sep=""),
            "DISPLAY_DEGREE"=3)
```

```
model <- crs(...,opts=opts)
```

will print out in gory detail exactly what the search engine is doing.

However, if this reveals that there something odd going on (i.e. you are seeing a lot of `inf` function values being printed out), then you might wish to begin by restricting the dimension of the combinatoric search process. By default `degree.max=10` for each predictor and `segments.max=10` as well. So this can lead to a basis with 21 columns for one predictor and when using `basis="tensor"` or `basis="auto"` (which computes both the tensor and additive bases) the dimension of the basis can swamp the number of observations in the sample (e.g. with 4 regressors we can have a tensor product multivariate basis that has up to $21^4=194481$ columns when using the default `degree.max=10` for each predictor and `segments.max=10`). For this illustration, the cross-validation function can approach ∞ if the sample size approaches 194491 from above and the search process will be searching for a non- ∞ value in order to proceed or terminate.

So, in such cases begin by restricting the dimension of the spline basis matrix by setting, for instance, `degree.max=2` and `segments.max=2`. Or begin by searching only over the spline degree by setting `complexity="degree"` (the default is `complexity="degree-knots"`). The routine will throw a warning if you have a solution that hits the maximum value of `degree.max` or `segments.max` and offer some practical advice in these cases.

Alternatively, restrict attention to additive (semiparametric) splines by setting `basis="additive"` (e.g. with 4 regressors we can have a tensor product multivariate basis that has up to $21 \times 4 = 84$ columns when using the default `degree.max=10` for each predictor and `segments.max=10`) at the cost of imposing additivity which can be restrictive.

Alternatively, consider kernel regression that does not suffer from this computational limitation (see e.g. the `np` package).

2.12. My estimated model is not ‘smooth’ (e.g. cross-validation chooses, say, the spline degree=1 and number of segments=3). How can I modify this? Unlike, say, smoothing splines that penalize ‘roughness’ (the second derivative of the estimate) and fix the spline degree at, say, three, regression splines fitted by cross-validation can deliver a model that minimizes the objective function without regard to smoothness (a strength, not a weakness in my opinion).

- (1) **All Predictors Relevant:** If, however, you wish an estimate that is, say, twice continuously differentiable, simply set `degree.min=3` (one degree higher than the desired degree of smoothness) and then determine the appropriate model via cross-validation subject to this constraint (i.e. re-estimate your model with the option `degree.min=3` for example). Otherwise, you can simply override what cross-validation delivered via `cv="none"` and `degree=c(3,3,...)` etc. (i.e. set the degree and segments in an ad-hoc fashion, which I would not recommend). Of course, this will not be optimal according to the cross-validation criterion but will achieve the desired degree of smoothness.
- (2) **Some Predictors Irrelevant:** In some cases the optimal spline degree will be zero or bandwidth one or inclusion indicator zero (the latter occurring whe `kernel=FALSE` is specified). This means that cross-validation has automatically removed one or more variables.

In this case first remove these variables from the model to be estimated, then re-run your model on the subset of relevant variables only but following the advice in (1) which will achieve the desired degree of smoothness subject to the irrelevant variables no longer being present in the final model.

2.13. I estimated a parametric model using the `lm()` function. How can I compare the cross-validation score from the `crs()` approach with that for the parametric model? This can be readily achieved for the parametric model as follows:

```
data(wage1)
model.lm <- lm(lwage ~ married + female + nonwhite + educ +
```

```

      exper + tenure, data = wage1)
cv.lm <- mean(residuals(model.lm)^2/(1-hatvalues(model.lm))^2)
cv.lm

```

You can then compare this with that for the `crs` model. If the cross-validation score is lower for one model, that indicates that the model possessing the lowest score is to be preferred.

2.14. Why do some runs result in a function value of 1.340781e+154 when conducting multistarting? As of version 0.15-1 we conduct extensive testing for ill-conditioned bases (univariate and multivariate) and adjust search limits accordingly. However, when a multivariate basis is ill-conditioned we apply a large penalty (`sqrt(.Machine$double.xmax)` which equals 1.340781e+154 on most processors). Though the search process will try to detect a minimum it can fail here if the objective function is ‘flat’ in a neighborhood of the initial values.

When this occurs you can either increase `nmulti` and/or decrease `degree.max` and restart the search.

Note also that as of version 0.15-1, the initial search values will be degree one and segment one (i.e. a linear model) unless you provide the vectors `degree=c(...)` and `segments=c(...)` which will then be used instead as the starting values for the first multi-start.

2.15. `snomadr` appears to be crashing. If you receive the message

Calling NOMAD (Nonsmooth Optimization by Mesh Adaptive Direct Search)

```

*** caught segfault ***
address 0x68, cause 'memory not mapped'

```

Traceback:

```

1: .Call(smulinomadRSolve, ret)

```

kindly first ensure that you have write privileges in your current directory (`snomadr` creates temporary files in the current working directory and if this operation fails you may receive this error).

2.16. How can I save a PDF of a plot created with the option `persp.rgl=TRUE`?

Version 0.15-1 has added support for RGL via the `rgl` package which is a 3D real-time rendering device driver system for R using OpenGL. These plots are dynamic so you can spin them and resize them using your keypad/mouse. However, they are not standard graphics objects that can be saved using R commands such as `pdf()`. But they can be saved as a

PDF by first calling `rgl` and then issuing the command `rgl.postscript("foo.pdf","pdf")` where `foo.pdf` is the desired name of your PDF file as the following illustrates:

```
n <- 1000
x1 <- sort(rnorm(n))
x2 <- rnorm(n)
y <- x1^3 + rnorm(n,sd=.1)
model <- crs(y~x1+x2)
plot(model,mean=T,persp.rgl=T)
rgl.postscript("foo.pdf","pdf")
```

However, this pdf driver does not support some features such as transparency etc. A better alternative is to create a `png` file as follows:

```
n <- 1000
x1 <- sort(rnorm(n))
x2 <- rnorm(n)
y <- x1^3 + rnorm(n,sd=.1)
model <- crs(y~x1+x2)
plot(model,mean=T,persp.rgl=T)
rgl.snapshot("foo.png")
```

and then include this in your \LaTeX document using `\includegraphics[scale=.5]{foo.png}`.

2.17. I have noticed that as more categorical predictors are added or as the number of outcomes for each categorical predictor increases, computation time increases when `kernel=TRUE`. Why is this so and can anything be done? When `kernel=TRUE` we need to compute kernel weighted regression for each unique combination of the categorical predictors. So if you have one binary predictor, estimation involves two calls to the weighted least squares solver. Now suppose that you have three categorical predictors each having $c = 10$ outcomes. Now estimation involves $10^3 = 1000$ calls to the weighted least squares solver. This will affect all aspects of estimation (cross-validation etc.) which results in increases in computation time.

As of version 0.15-8 and up, you can potentially reduce computation time *when you have categorical predictors* by discretizing the bandwidth `lambda` (when `lambda.discrete=TRUE` we divide `lambda` $\in [0,1]$ into `lambda.discrete.num=100 (+1)` values, $(0/100, 1/100, \dots, 100/100)$) so that integer (discrete) search via NOMAD will be undertaken rather than continuous search which is done by setting `lambda.discrete=TRUE`.

This can potentially reduce run-time with little expected loss in accuracy for the cross-validated search procedure. However, when there are many categorical predictors each having > 2 outcomes, search based on `lambda.discrete=TRUE` may be more likely to become ensnared in local minima than the default (i.e. properly treating the bandwidths for the categorical predictors as continuous).

To reduce computation time further, you can also decrease the number of times search is restarted from different (random) starting points by setting `nmulti=i` to $i < 5$ (the default is 5) but I would caution against doing this except when conducting exploratory data analysis (the objective function is typically nonsmooth and may possess multiple local minima).

REFERENCES

- [1] Francisco Cribari-Neto and Spyros G Zarkos. R: Yet another econometric programming environment. *Journal of Applied Econometrics*, 14(3):319–29, May-June 1999. Available at <http://ideas.repec.org/a/jae/japmet/v14y1999i3p319-29.html>.
- [2] Yves Croissant. *Ecdat: Data sets for econometrics*, 2011. R package version 0.1-6.1.
- [3] Grant V. Farnsworth. Econometrics in R. Technical report, October 2008. Available at <http://cran.r-project.org/doc/contrib/Farnsworth-EconometricsInR.pdf>.
- [4] Christian Kleiber and Achim Zeileis. *Applied Econometrics with R*. Springer-Verlag, New York, 2008. ISBN 978-0-387-77316-2.
- [5] J. S. Racine and R. Hyndman. Using R to teach econometrics. *Journal of Applied Econometrics*, 17(2):175–189, 2002.

CHANGES FROM VERSION 0.15-9 TO 0.15-10 [3-DEC-2011]

- Fixed regression in code when using `cv.func=cv.aic`
- Added option `bwtype="auto"` for `npglmreg` (automatically determine the bandwidth type via cross-validation)

CHANGES FROM VERSION 0.15-8 TO 0.15-9 [25-NOV-2011]

- Fixed glitch when computing/plotting derivatives with `crs` and multiple predictors are of degree zero
- Added support for parallel `npglmreg` (calls `npRmpi` rather than `np` - see demo)
- Improved handling of complex bases via `dim.bs` (test for negative degrees of freedom prior to attempting to construct the basis function - dramatically reduces memory overhead and can cut down on unnecessary computation)

CHANGES FROM VERSION 0.15-7 TO 0.15-8 [16-NOV-2011]

- Added support for `is.fullrank` testing for generalized local polynomial kernel regression for cross-validation (default remains ridging a la Seifert & Gasser), replaced `rcond(t(X)%*%X)` with `is.fullrank(X)` (smaller memory footprint for large datasets)
- Fixed glitch with derivative computation when one or more `degree` is 0 when `kernel=TRUE` and additionally when one or more factors is excluded when using `kernel=FALSE`
- Fixed issue with reported cross-validation score only corresponding to leave-one-out cross-validation by passing back `cv` function from solver rather than computing post estimation
- Added ability to estimate quantile regression splines
- More rigorous testing for rank deficient fit via `rcond()` in `cv` function
- Fixed issue where `degree.min` was set > 1 but initial degree was 1 (corresponding to the linear model which is the default for the initial degree otherwise)
- Added the option to treat the continuous bandwidths as discrete with `lambda.discrete.num+1` values in the range $[0,1]$ which can be more computationally efficient when a ‘quick and dirty’ solution is sufficient rather than conducting mixed integer search treating the `lambda` as real-valued
- Corrected incorrect warning about using `basis="auto"` when there was only one continuous predictor

CHANGES FROM VERSION 0.15-6 TO 0.15-7 [24-OCT-2011]

- Added logical `model.return` to `crs` (default `model.return=FALSE`) which previously returned a list of models corresponding to each unique combination of the categorical predictors when `kernel=TRUE` (the memory footprint could be potentially very large so this allows the user to generate this list if so needed)

CHANGES FROM VERSION 0.15-5 TO 0.15-6 [17-OCT-2011]

- Compiler error thrown on some systems due to changes in `Eval_Point.hpp` corrected

CHANGES FROM VERSION 0.15-4 TO 0.15-5 [16-OCT-2011]

- Thanks to Professor Brian Ripley, additional Solaris C/C++ compiler warnings/issues have been resolved
- Some internal changes for soon to be deprecated functionality (`sd(<matrix>)`) expected to be deprecated shortly)

CHANGES FROM VERSION 0.15-3 TO 0.15-4 [15-OCT-2011]

- Extended the `gsl.bs` functionality to permit out-of-sample prediction of the spline basis and its derivatives
- Added option `knots="auto"` to automatically determine via cross-validation whether to use quantile or uniform knots
- Minor changes to help page examples and descriptions and to the `crs` vignette

CHANGES FROM VERSION 0.15-2 TO 0.15-3 [05-SEPT-2011]

- Fixed glitch when all degrees are zero when computing the cross-validation function (also fixes glitch when all degrees are zero when plotting the partial surfaces)
- Added new function `crssigtest` (to be considered in beta status until further notice)
- Added F test for no effect (joint test of significance) in `crs summary`
- Both degree and segments now set to one for first multistart in `crs` (previously only degree was, but intent was always to begin from a linear model (with interactions where appropriate) so this glitch is corrected)
- Test for pathological case in `npplpreg` when initializing bandwidths where `IQR` is zero but `sd > 0` (for setting robust `sd`) which occurs when there exist many repeated values for a continuous predictor
- Added ‘typical usage’ preformatted illustrations for docs

CHANGES FROM VERSION 0.15-1 TO 0.15-2 [30-JULY-2011]

- Renamed COPYING file to COPYRIGHTS

CHANGES FROM VERSION 0.15-0 TO 0.15-1 [29-JUL-2011]

- Automated detection of ordered/unordered factors implemented
- Initial degree values set to 1 when conducting NOMAD search (only for initial, when `nmulti > 1` random valid values are generated)
- Multiple tests for well-conditioned B-spline bases, dynamic modification of search boundaries when ill-conditioned bases are detected, and detection of non-positive degrees of freedom and full column rank of the spline basis (otherwise the penalty `sqrt(.Machine$double.xmax)` is returned during search) - this can lead to a significant reduction in the memory footprint
- Added support for generalized B-spline kernel bases (varying order generalized polynomial)
- Corrected issue with plot when variables were cast as `factor` in the model formula
- Fixed glitch with return object and i/o when `cv="bandwidth"` and `degree=c(0,0,...,0)`
- Added tests for pathological cases (e.g. optimize degree and knots but set max degree to min degree or max segments to min hence no search possible).
- Added argument `cv.threshold` that uses exhaustive search for simple cases where the number of objective function evaluations is less than `cv.threshold` (currently set to 1000 but user can set). Naturally exhaustive search is always preferred but often unfeasible, so when it is feasible use it.
- Added additional demos for constrained estimation (Du, Parmeter, and Racine (2011)), inference, and a sine-based function.
- Substantial reductions in run-time realized.
 - Product kernel computation modified for improved run-time of kernel-based cross-validation and estimation.
 - Moved from `lsfit` to `lm.fit` and from `lm` to `lm.wfit/lm.fit` in `cv.kernel.spline` and `cv.factor.spline` (compute objective function values). Two effects - R devel indicates `lm.fit/lm.wfit` are more robust (confirmed for large number of predictors) and much faster `cv.kernel.spline` function emerges (run-time cut 20-30%).
 - The combined effects of these changes are noticeable. For instance, run-time for `wage1` with 7 predictor cross-validation goes from 510 seconds in 0.15-0 to 304

seconds due to use of `lm.fit/lm.wfit` described below to 148 seconds due to the modified kernel function.

CHANGES FROM VERSION 0.14-9 TO 0.15-0 [23-JUN-2011]

- Thanks to Professor Brian Ripley, compile on Solaris system issues are resolved, and check/examples are reduced in run time to alleviate the excessive check times by the R development team. Many thanks to them for their patience and guidance.
- Minor changes to `radial_rgl` demo

CHANGES FROM VERSION 0.14-8 TO 0.14-9 [20-JUN-2011]

- Cleaned up issues for creating binary for windows
- Setting seed in `snomadr.cpp` via `snomadr.R` for starting points when `nmulti > 0`
- Increased default `MAX_BB_EVAL` from 500 to 10000 (makes a difference for difficult problems) and modified default `EPSILON` in `NOMAD` along with other parameters (`MIN_MESH_SIZE`, `MIN_POLL_SIZE`) to reflect actual machine precision (using R's `.Machine$double.eps` where `NOMAD` fixed `EPSILON` at `1e-13`)
- Zhenghua added help functionality for retrieving help via `snomadr`
- Now default number of restarts in `crs` is 5 (zero is not reliable and I want sensible defaults in this package - higher is better but for many problems this ought to suffice)
- Corrected glitches in interactive demos where options were not being passed, updated docs to reflect demos

CHANGES FROM VERSION 0.14-7 TO 0.14-8 [10-JUN-2011]

- `crsiv` now returns a `crs` model object that supports residuals, fitted, predict and other generic functions. Note that this approach is based on first computing the model via regularization and then feeding a transformed response to a `crs` model object. You can test how close the two approaches are to one another by comparing `model$phi_hat` with `fitted(model)` via

```
all.equal(as.numeric(fitted(model)), as.numeric(model$phi_hat))
```

VERSION 0.14-7 [09-JUN-2011]

- Initial release of the `crs` package on CRAN.